


Lectures on Natural Language Processing

12. Conditional Random Fields

Karl Stratos

Review: Hidden Markov Model (HMM)

A generative model of sequence labeling




$$p(x_1 \dots x_T, y_1 \dots y_T) = \prod_{t=1}^T \underbrace{\tau(y_t | y_{t-1})}_{\text{transition prob}} \times \underbrace{o(x_t | y_t)}_{\text{emission prob}} \times \tau(y_* | y_T)$$

- ▶ **Forward/backward algorithm:** Marginalization in $O(T |\mathcal{Y}|^2)$
- ▶ **Viterbi decoding:** Best label sequence in $O(T |\mathcal{Y}|^2)$
- ▶ **Max marginal decoding:** Best label per position in $O(T |\mathcal{Y}|^2)$

Conditional Random Field (CRF) Tagger

A *discriminative* model of sequence labeling



$$p_{\theta}(y_1 \dots y_T | x_1 \dots x_T) = \frac{\prod_{t=1}^T \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, y_{t-1}, y_t, t))}{\sum_{y'_1 \dots y'_T \in \mathcal{Y}} \prod_{t=1}^T \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, y'_{t-1}, y'_t, t))}$$

Equivalently a “giant softmax” over label sequences where the score function is first-order Markovian

$$\mathbf{score}_{\theta}(x_1 \dots x_T, y_1 \dots y_T) = \sum_{t=1}^T \mathbf{score}_{\theta}(\underbrace{x_1 \dots x_T}_{\text{all inputs}}, \underbrace{y_{t-1}}_{\text{only prev. label}}, y_t, t)$$

Neural Parameterization of CRF Taggers

- ▶ $\text{score}_\theta(x_1 \dots x_T, y, y', t) \in \mathbb{R}$ should capture how likely the t -th label is y' , given $x_1 \dots x_T$ and the previous label y
- ▶ Example: Assuming some contextual embeddings $\text{enc}_\theta(x_1 \dots x_T) \in \mathbb{R}^{T \times d}$ (e.g., BiLSTM, BERT)


$$O = \text{enc}_\theta(x_1 \dots x_T)W \in \mathbb{R}^{T \times |\mathcal{Y}|}$$

$$\text{score}_\theta(x_1 \dots x_T, \textcolor{blue}{y}, \textcolor{red}{y}', t) = T_{\textcolor{blue}{y}, \textcolor{red}{y}'} + O_{t, \textcolor{red}{y}'}$$

- ▶ Learnable parameters
 - ▶ $W \in \mathbb{R}^{d \times |\mathcal{Y}|}$ computes per-position label logits (i.e., O captures the “emission” scores)
 - ▶ $T \in \mathbb{R}^{|\mathcal{Y}| \times |\mathcal{Y}|}$ captures the “transition” scores
- ▶ Flexible, e.g., could define transition score $y \rightarrow y'$ to be $v_y^\top A v_{y'}$ where $v_y \in \mathbb{R}^{d'}$ is a learnable embedding of label y

Another View of CRF: Product of Potential Functions

Each maximal clique (i.e., fully connected subgraph) is associated with a nonnegative “potential function”



$$\psi_{\theta}^t(x_1 \dots x_T, y, y') = \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, y, y', t)) \geq 0$$

CRF distribution = normalized product of potential functions

$$p_{\theta}(y_1 \dots y_T | x_1 \dots x_T) \propto \prod_{t=1}^T \psi_{\theta}^t(x_1 \dots x_T, y_{t-1}, y_t)$$

Implicit Backward Sampling in CRF Tagger

Lemma: The CRF tagger implies a first-order backward sampling procedure.

$$y_1 \dots y_T \sim p_\theta(\cdot | x_1 \dots x_T)$$

|||

$$y_T \sim q_\theta(\cdot | x_1 \dots x_T, T)$$

$$y_{T-1} \sim q_\theta(\cdot | x_1 \dots x_T, y_T, T-1)$$

$$y_{T-2} \sim q_\theta(\cdot | x_1 \dots x_T, y_{T-1}, T-2)$$

\vdots

$$y_1 \sim q_\theta(\cdot | x_1 \dots x_T, y_2, 1)$$

This fact allows us to efficiently sample a label sequence from the CRF tagger.

The Marginalization and Inference Problems

How can we compute the normalizer, aka. **partition function**?

$$Z_{\theta}(x_1 \dots x_T) = \sum_{y_1 \dots y_T \in \mathcal{Y}} \prod_{t=1}^T \psi_{\theta}^t(x_1 \dots x_T, y_{t-1}, y_t)$$

How can we find the most likely label sequence?

$$y_1^* \dots y_T^* = \arg \max_{y_1 \dots y_T \in \mathcal{Y}} \prod_{t=1}^T \psi_{\theta}^t(x_1 \dots x_T, y_{t-1}, y_t)$$

(Normalization not necessary for inference)

Forward Algorithm for CRFs

Dynamic programming: Given $x_1 \dots x_T$, we fill out a table $\alpha \in \mathbb{R}^{T \times |\mathcal{Y}|}$ left-to-right where

$$\alpha(t, y) = \sum_{y_1 \dots y_t \in \mathcal{Y}: y_t = y} \prod_{l=1}^t \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l)$$

Base case?

$$\alpha(1, y) =$$

Forward Algorithm for CRFs

Dynamic programming: Given $x_1 \dots x_T$, we fill out a table $\alpha \in \mathbb{R}^{T \times |\mathcal{Y}|}$ left-to-right where

$$\alpha(t, y) = \sum_{y_1 \dots y_t \in \mathcal{Y}: y_t = y} \prod_{l=1}^t \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l)$$

Base case?

$$\alpha(1, y) = \psi_{\theta}^1(x_1 \dots x_T, y_0, y)$$

Forward Algorithm for CRFs: Main Body ($t > 1$)

$$\begin{aligned}\alpha(t, y') &= \sum_{y_1 \dots y_t: y_t = y'} \prod_{l=1}^t \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l) \\&= \sum_{y_1 \dots y_{t-1}} \left(\prod_{l=1}^{t-1} \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l) \right) \psi_{\theta}^t(x_1 \dots x_T, y_{t-1}, y') \\&= \sum_{\mathbf{y}} \sum_{y_1 \dots y_{t-1}: y_{t-1} = \mathbf{y}} \left(\prod_{l=1}^{t-1} \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l) \right) \psi_{\theta}^t(x_1 \dots x_T, \mathbf{y}, y') \\&= \sum_{\mathbf{y}} \alpha(t-1, \mathbf{y}) \psi_{\theta}^t(x_1 \dots x_T, \mathbf{y}, y')\end{aligned}$$

Viterbi Algorithm for CRFs

Given $x_1 \dots x_T$, we fill out a table $\pi \in \mathbb{R}^{T \times |\mathcal{Y}|}$ left-to-right where

$$\pi(t, y) = \max_{y_1 \dots y_t \in \mathcal{Y}: y_t = y} \prod_{l=1}^t \psi_{\theta}^l(x_1 \dots x_T, y_{l-1}, y_l)$$

As in HMMs, same as forward except we switch sum with max!

$$\pi(1, y) = \psi_{\theta}^1(x_1 \dots x_T, y_0, y)$$

$$\pi(t, y') = \max_y \pi(t-1, y) \psi_{\theta}^t(x_1 \dots x_T, y, y')$$


$$\mathbf{bp}(t, y') = \arg \max_y \pi(t-1, y) \psi_{\theta}^t(x_1 \dots x_T, y, y')$$

Extract the argmax label sequence by backtracking

$$y_T^* = \arg \max_{y \in \mathcal{Y}} \pi(T, y) \quad y_{t-1}^* = \mathbf{bp}(t, y_t^*) \quad \text{for } t = T \dots 2$$

Training a CRF Tagger

Given labeled sequences, calculate cross-entropy in $O(T|\mathcal{Y}|^2)$ using the forward algorithm (in log space for numerical stability):



$$\max_{\theta} \sum_{i=1}^N \log p_{\theta}(y_1^{(i)} \dots y_{T_i}^{(i)} | x_1^{(i)} \dots x_{T_i}^{(i)})$$


(Backprop "through structure")

After training, infer tags for any $x_1 \dots x_T$ in $O(T|\mathcal{Y}|^2)$ using the Viterbi algorithm:

$$y_1^* \dots y_T^* = \arg \max_{y_1 \dots y_T \in \mathcal{Y}} \log p_{\theta}(y_1 \dots y_T | x_1 \dots x_T)$$


Markov Random Fields (MRFs)

- ▶ CRF tagger is a special case of a **Markov random field** (MRF) (= *undirected* graphical model)
- ▶ MRF: Joint distribution that factorizes over *maximal* cliques C equipped with nonnegative **potential functions** ψ_C
 - ▶ Clique: Fully connected subgraph
 - ▶ Maximal clique: Clique that loses full connectivity if any node is added



$$\Pr(X_{1:6}) = \frac{1}{Z} \underbrace{\psi_{1:3}(X_{1:3})}_{\geq 0} \underbrace{\psi_{3:6}(X_{3:6})}_{\geq 0}$$
$$Z = \sum_{x_{1:6}} \psi_{1:3}(x_{1:3}) \psi_{3:6}(x_{3:6})$$


- ▶ Factor graph notation



$$\Pr(X_{1:6}) \propto \psi_{1:3}(X_{1:3}) \psi_{3:6}(X_{3:6})$$

General Marginalization and Inference in MRFs


- Some variables are observed in an MRF, work with conditional distributions (i.e., “conditional” random fields):



$$\Pr(X_{1:2}, X_{4:6} | X_3 = c) = \frac{1}{Z(X_3 = c)} \psi_{1:3}(X_1 X_2 c) \psi_{3:6}(X_3 c)$$
$$Z(X_3 = c) = \sum_{x_{1:6}: x_3 = c} \psi_{1:3}(x_{1:3}) \psi_{3:6}(x_{3:6})$$

- MRF again poses general structured prediction problems, like
 - Marginalize: $\Pr(X_5 = c' | X_3 = c)$
 - Infer: $\arg \max_{x_{1:2}, x_{4:6}} \Pr(X_{1:2} = x_{1:2}, X_{4:6} = x_{4:6} | X_3 = c)$
- Variable elimination (VE).** General “recipe” to solve these problems exactly in $O(n_{\text{infer}} K^{C_{\text{max}}})$ time (assuming no cycles) where
 - n_{infer} : Number of variables in MRF that we’re inferring
 - K : Number of possible values that variables can take
 - C_{max} : Size of the *largest* maximal clique (e.g., 2 in taggers)
- Too abstract to be directly useful (e.g., must specify elimination ordering), but provides a unified framework of structured prediction (e.g., forward, Viterbi are VE on chains)

Example MRF




$$p(x_1, x_2, x_3, x_4, x_5) = \frac{\psi_{12}(x_1, x_2)\psi_{13}(x_1, x_3)\psi_{14}(x_1, x_4)\psi_{25}(x_2, x_5)\psi_{35}(x_3, x_5)\psi_{45}(x_4, x_5)}{\sum_{x' \in \{1 \dots K\}^5} \psi_{12}(x'_1, x'_2)\psi_{13}(x'_1, x'_3)\psi_{14}(x'_1, x'_4)\psi_{25}(x'_2, x'_5)\psi_{35}(x'_3, x'_5)\psi_{45}(x'_4, x'_5)}$$

Computing the normalizer Z naively will take $O(K^5)$ time.


Example MRF: Variable Elimination

VE ordering X_2, X_3, X_4, X_5, X_1 :



Example MRF: Variable Elimination


VE ordering X_2, X_3, X_4, X_5, X_1 :



$$Z = \sum_{x_3, x_4, x_5, x_1} \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{12}(x_1, x_2) \psi_{25}(x_2, x_5) \right)}_{\phi_2(x_1, x_5)}$$

Example MRF: Variable Elimination


VE ordering X_2, X_3, X_4, X_5, X_1 :



$$\begin{aligned} Z &= \sum_{x_3, x_4, x_5, x_1} \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{12}(x_1, x_2) \psi_{25}(x_2, x_5) \right)}_{\phi_2(x_1, x_5)} \\ &= \sum_{x_4, x_5, x_1} \psi_{14}(x_1, x_4) \psi_{45}(x_4, x_5) \phi^2(x_1, x_5) \underbrace{\left(\sum_{x_3} \psi_{13}(x_1, x_3) \psi_{35}(x_3, x_5) \right)}_{\phi_3(x_1, x_5)} \end{aligned}$$

Example MRF: Variable Elimination


VE ordering X_2, X_3, X_4, X_5, X_1 :



$$\begin{aligned} Z &= \sum_{x_3, x_4, x_5, x_1} \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{12}(x_1, x_2) \psi_{25}(x_2, x_5) \right)}_{\phi_2(x_1, x_5)} \\ &= \sum_{x_4, x_5, x_1} \psi_{14}(x_1, x_4) \psi_{45}(x_4, x_5) \phi^2(x_1, x_5) \underbrace{\left(\sum_{x_3} \psi_{13}(x_1, x_3) \psi_{35}(x_3, x_5) \right)}_{\phi_3(x_1, x_5)} \\ &= \sum_{x_5, x_1} \phi^2(x_1, x_5) \phi^3(x_1, x_5) \underbrace{\left(\sum_{x_4} \psi_{14}(x_1, x_4) \psi_{45}(x_4, x_5) \right)}_{\phi_4(x_1, x_5)} \end{aligned}$$

Example MRF: Variable Elimination

VE ordering X_2, X_3, X_4, X_5, X_1 :




$$\begin{aligned} Z &= \sum_{x_3, x_4, x_5, x_1} \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{12}(x_1, x_2) \psi_{25}(x_2, x_5) \right)}_{\phi_2(x_1, x_5)} \\ &= \sum_{x_4, x_5, x_1} \psi_{14}(x_1, x_4) \psi_{45}(x_4, x_5) \phi^2(x_1, x_5) \underbrace{\left(\sum_{x_3} \psi_{13}(x_1, x_3) \psi_{35}(x_3, x_5) \right)}_{\phi_3(x_1, x_5)} \\ &= \sum_{x_5, x_1} \phi^2(x_1, x_5) \phi^3(x_1, x_5) \underbrace{\left(\sum_{x_4} \psi_{14}(x_1, x_4) \psi_{45}(x_4, x_5) \right)}_{\phi_4(x_1, x_5)} \end{aligned}$$

Runtime: $O(K^3)$


Example MRF: Variable Elimination, Different Ordering

VE ordering X_1, X_2, X_3, X_4, X_5 :



Example MRF: Variable Elimination, Different Ordering


VE ordering X_1, X_2, X_3, X_4, X_5 :



$$Z = \sum_{x_2, x_3, x_4, x_5} \psi_{25}(x_2, x_5) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_1} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \right)}_{\psi^1(x_2, x_3, x_4)}$$

Example MRF: Variable Elimination, Different Ordering


VE ordering X_1, X_2, X_3, X_4, X_5 :



$$\begin{aligned} Z &= \sum_{x_2, x_3, x_4, x_5} \psi_{25}(x_2, x_5) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_1} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \right)}_{\psi^1(x_2, x_3, x_4)} \\ &= \sum_{x_3, x_4, x_5} \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{25}(x_2, x_5) \psi^1(x_2, x_3, x_4) \right)}_{\psi^2(x_3, x_4, x_5)} \end{aligned}$$

Example MRF: Variable Elimination, Different Ordering


VE ordering X_1, X_2, X_3, X_4, X_5 :



$$\begin{aligned}
 Z &= \sum_{x_2, x_3, x_4, x_5} \psi_{25}(x_2, x_5) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_1} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \right)}_{\psi^1(x_2, x_3, x_4)} \\
 &= \sum_{x_3, x_4, x_5} \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{25}(x_2, x_5) \psi^1(x_2, x_3, x_4) \right)}_{\psi^2(x_3, x_4, x_5)} \\
 &= \sum_{x_4, x_5} \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_3} \psi_{35}(x_3, x_5) \psi^2(x_3, x_4, x_5) \right)}_{\psi^3(x_4, x_5)}
 \end{aligned}$$

Example MRF: Variable Elimination, Different Ordering

VE ordering X_1, X_2, X_3, X_4, X_5 :




$$\begin{aligned} Z &= \sum_{x_2, x_3, x_4, x_5} \psi_{25}(x_2, x_5) \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_1} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{14}(x_1, x_4) \right)}_{\psi^1(x_2, x_3, x_4)} \\ &= \sum_{x_3, x_4, x_5} \psi_{35}(x_3, x_5) \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_2} \psi_{25}(x_2, x_5) \psi^1(x_2, x_3, x_4) \right)}_{\psi^2(x_3, x_4, x_5)} \\ &= \sum_{x_4, x_5} \psi_{45}(x_4, x_5) \underbrace{\left(\sum_{x_3} \psi_{35}(x_3, x_5) \psi^2(x_3, x_4, x_5) \right)}_{\psi^3(x_4, x_5)} \end{aligned}$$

Runtime: $O(K^4)$

General Tagging with MRFs


- ▶ No independence assumptions: $O(T |\mathcal{Y}|^T)$



$$p_{\theta}(y_{1:3}|x_{1:3}) \propto \exp(\mathbf{score}_{\theta}(x_{1:3}, y_{1:3}))$$

$$C_{\max} = 3$$


- ▶ Greedy tagging (i.e., softmax per position): $O(T |\mathcal{Y}|)$



$$p_{\theta}(y_{1:3}|x_{1:3}) \propto \prod_{t=1}^3 \exp(\mathbf{score}_{\theta}(x_{1:3}, y_t, t))$$

$$C_{\max} = 1$$

- ▶ First-order CRF: $O(T |\mathcal{Y}|^2)$




$$p_{\theta}(y_{1:3}|x_{1:3}) \propto \prod_{t=1}^3 \exp(\mathbf{score}_{\theta}(x_{1:3}, y_{t-1}, y_t, t))$$

$$C_{\max} = 2$$

More Facts About Graphical Models


Any directed (acyclic) graph can be expressed by an equivalent MRF



- ▶ Forward algorithm for HMM: VE with left-to-right elimination ordering, generalizable to trees

VE applicable only if there's no cycle (e.g., sequences, trees)


- ▶ If cycle between unobserved variables, $O(n_{\text{infer}} K^{C_{\text{max}}})$ runtime guarantee doesn't hold, e.g., marginalization intractable in



- ▶ Can technically combine factors until there's no cycle and apply VE, but that's no better than brute-force
- ▶ Efficient approximations possible: loopy belief propagation

Review: PCFG

A top-down generative model of parsing

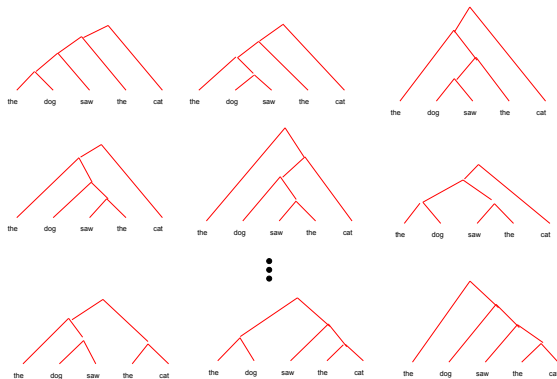


$$\begin{aligned} p(\text{tree}) &= q(S \rightarrow \text{NP VP}) \times q(\text{NP} \rightarrow \text{D N}) \\ &\quad \times q(\text{D} \rightarrow \text{the}) \times q(\text{N} \rightarrow \text{dog}) \\ &\quad \times q(\text{VP} \rightarrow \text{V P}) \times q(\text{V} \rightarrow \text{saw}) \times q(\text{P} \rightarrow \text{him}) \end{aligned}$$

- ▶ **Inside algorithm:** Marginalization in $O(T^3 |\mathcal{R}|)$
- ▶ **CKY decoding:** Best tree in $O(T^3 |\mathcal{R}|)$
- ▶ **Max marginal decoding:** Outside algorithm $O(T^3 |\mathcal{R}|)$, labeled recall algorithm $O(T^3)$

Discriminative Parsing

Given a sentence, define a conditional distribution over all possible (binary) trees




Questions:

1. How do we represent a tree?
2. How do we define the score of a tree to allow for efficient calculations?

CRF Parser

Unlabeled binary tree = a valid set of spans



Conditional tree probability as normalized product of potentials

$$p_{\theta}(\tau|x_1 \dots x_T) \propto \prod_{(s,t) \in \tau} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, \mathbf{s}, \mathbf{t}))$$

Equivalently a “giant softmax” over trees, where the tree score is

$$\mathbf{score}_{\theta}(x_1 \dots x_T, \tau) = \sum_{(s,t) \in \tau} \mathbf{score}_{\theta}(x_1 \dots x_T, \mathbf{s}, \mathbf{t})$$

Neural Parameterization of CRF Parsers

- $\text{score}_\theta(x_1 \dots x_T, i, j) \in \mathbb{R}$ should capture how likely (i, j) is a span in the underlying tree

- Example: Given some text encoder (e.g., BERT)


$$h_1 \dots h_T = \text{enc}_\theta(x_1 \dots x_T)$$

$$\text{score}_\theta(x_1 \dots x_T, i, j) = v^\top \text{ReLU}(W(h_j - h_i) + b) + b'$$

- Can also use labeled score (e.g., for constituency parsing)

$$\text{score}_\theta(x_1 \dots x_T, i, j, l) = v_l^\top \text{ReLU}(W(h_j - h_i) + b) + b'_l$$

- Given training data $(x^{(1)}, \tau^{(1)}) \dots (x^{(N)}, \tau^{(N)})$ a CRF parser can be trained by cross-entropy or max-margin loss




IF we can do marginalization/inference!

(Image credit: Kitaev and Klein (2018))

Notation: Set of All Trees Over Some Span

$B_T[i, j]$: all binary trees over the span $[i, j]$ in a sequence of length T



$$B_5[2, 4] = \{ \{ [2, 4], [2, 3] \}, \{ [2, 4], [3, 4] \} \}$$

Span-conditional tree distribution by

$$p_{\theta}(\tau | x_1 \dots x_T, i, j) = \frac{\exp(\mathbf{score}_{\theta}(x_1 \dots x_T, \tau))}{\sum_{\tau' \in B_T[i, j]} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, \tau'))} \quad \forall \tau \in B_T[i, j]$$

Special case: final CRF parser


$$p_{\theta}(\tau | x_1 \dots x_T) = p_{\theta}(\tau | x_1 \dots x_T, 1, T)$$

Implicit Top-Down Sampling in CRF Parser

Lemma: The CRF parser implies a top-down sampling procedure.

$$\tau \sim p_{\theta}(\cdot | x_1 \dots x_T, i, j)$$

|||



$$k \sim q_{\theta}(\cdot | x_1 \dots x_T, i, j)$$

$$\tau_{\text{left}} \sim p_{\theta}(\cdot | x_1 \dots x_T, i, k)$$

$$\tau_{\text{right}} \sim p_{\theta}(\cdot | x_1 \dots x_T, k+1, T)$$

Marginalization and Inference

Given a sentence $x_1 \dots x_T$,

1. How can we compute the partition function?

$$\sum_{\tau \in B_T[1,T]} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, \tau))$$

2. What is the most likely tree?

$$\tau^{\star} = \arg \max_{\tau \in B_T[1,T]} \mathbf{score}_{\theta}(x_1 \dots x_T, \tau)$$

Can use variants of the inside/CKY algorithm for PCFGs

Inside Algorithm for CRF Parser

Dynamic programming: Given $x_1 \dots x_T$, we fill out a table $\alpha \in \mathbb{R}^{T \times T}$ bottom up:


$$\alpha(i, j) = \sum_{\tau \in B_T[i, j]} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, \tau))$$

Base case?

$$\alpha(i, i) = \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, i, i))$$

Inside Algorithm for CRF Parser: Main Body

$$\begin{aligned}
 \alpha(i, j) &= \sum_{\tau \in B_T[i, j]} \prod_{(s, t) \in \tau} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, s, t)) \\
 &= \sum_{\substack{i \leq k < j \\ l \in B_T[i, k] \\ r \in B_T[k+1, j]}} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, i, j)) \times \left(\prod_{(s, t) \in l} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, s, t)) \right) \\
 &\quad \times \left(\prod_{(a, b) \in r} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, a, b)) \right) \\
 &= \sum_{i \leq k < j} \exp(\mathbf{score}_{\theta}(x_1 \dots x_T, i, j)) \times \alpha(i, k) \times \alpha(k+1, j)
 \end{aligned}$$



Introducing Latent Variables in Generative Models

- ▶ Generative models (e.g., LMs) define $p_{\theta}(x)$
 - ▶ The only random variable is observation x
- ▶ Idea: Introduce additional variable z and explicitly model an unseen generative process
 - ▶ We believe the process to be true (or at least useful for something), even though we don't observe it



(Original Image: [4edges/Wikimedia Commons](#))

Latent-Variable Generative Models (LVGMs)


- ▶ p_θ defining a *joint* distribution over observation $x \in \mathcal{X}$ and latent variable $z \in \mathcal{Z}$

$$p_\theta(x, z) = \underbrace{\kappa_\theta(x|z)}_{\text{conditional likelihood}} \times \underbrace{\pi_\theta(z)}_{\text{prior}}$$

- ▶ Very general definition
 - ▶ Can be discrete, continuous, or mixed
 - ▶ x can be structured, z can be structured, or both
- ▶ Why introduce latent variables?
 1. Clear generative story: Sample $z \sim \pi_\theta(z)$, then $x \sim \kappa_\theta(\cdot|z)$
 2. *Marginal* observation distribution can be more expressive
 3. Latent variables can be useful: Controllable generation (i.e., change z to get x we want), z natural representation of x

Marginal Observation Distribution

- ▶ LVGM defines a **marginal** distribution m_θ over \mathcal{X}
 - ▶ If z is discrete: $m_\theta(x) = \sum_{z \in \mathcal{Z}} p_\theta(x, z)$
 - ▶ If z is continuous: $m_\theta(x) = \int_{z \in \mathcal{Z}} p_\theta(x, z) dz$
 - ▶ If z is mixed: sum/integrate out appropriate dimensions
- ▶ m_θ can express a larger family of distributions
- ▶ Example: Bimodal distribution over $\mathcal{X} = \mathbb{R}$ cannot be expressed by any single Gaussian $\mathcal{N}(\mu, \sigma^2)$



- ▶ But can be expressed by a mixture of two Gaussians:

$$m_\theta(x) = \pi_1 \mathcal{N}(\mu_1, \sigma_1^2)(x) + \pi_2 \mathcal{N}(\mu_2, \sigma_2^2)(x)$$

Discrete latent variable $\mathcal{Z} = \{1, 2\}$

Better Explanation of Data

- ▶ Suppose iid samples from unknown **pop** over $\{a, b\}^{10}$ look like

$$x^{(1)} = (a, a, a, a, a, a, a, a, a, a) \quad x^{(2)} = (b, b, b, b, b, b, b, b, b, b)$$

$$x^{(3)} = (a, a, a, a, a, a, a, a, a, a) \quad x^{(4)} = (b, b, b, b, b, b, b, b, b, b)$$

$$x^{(5)} = (a, a, a, a, a, a, a, a, a, a) \quad x^{(6)} = (b, b, b, b, b, b, b, b, b, b)$$

- ▶ Bag-of-words model $p_{\theta}(x) = \prod_{j=1}^{10} p_{\theta}(x_j)$?
 - ▶ The model's independence assumption is clearly wrong!
 - ▶ Poor data fit: At most $p_{\theta}(x^{(i)}) = 2^{-10} < 0.001$ for each i
- ▶ LVGM $m_{\theta}(x) = \sum_{z \in \{1,2\}} \pi_{\theta}(z) \times \prod_{j=1}^{10} \kappa_{\theta}(x_j|z)$
 - ▶ The model makes the right assumption (draw a latent “topic” z and draw observation conditioned on z).
 - ▶ Can achieve $m_{\theta}(x^{(i)}) = 2^{-1}$ for each i with only twice more parameters
 - ▶ Also likely to generalize better (i.e., higher log likelihood of future samples)

Example LVGMs

- ▶ **HMMs:** $z \in \mathcal{Z}^T$ (unobserved label sequence), $x \in \mathcal{V}^T$ (sentence)

$$p_{\theta}(x, z) = \prod_{t=1}^{T+1} t_{\theta}(z_t | z_{t-1}) \times \prod_{t=1}^T o_{\theta}(x_t | z_t)$$

- ▶ **Gaussian LM:** $z \in \mathbb{R}^d$ (“thought vector”), $x \in \mathcal{V}^T$ (sentence)

$$p_{\theta}(x, z) = \mathcal{N}(0_d, I_{d \times d})(z) \times \prod_{t=1}^{T+1} p_{\theta}(x_t | x_{<t}, z)$$

- ▶ **Document hashing:** $z \in \{0, 1\}^d$ (“hash code”), $x \in \mathbb{R}^V$ (TFIDF document encoding)

$$p_{\theta}(x, z) = \prod_{j=1}^d \text{Bernoulli}(\lambda_j)(z_j) \times \prod_{k=1}^V p_{\theta}(x_k | z)$$

Marginal Log-Likelihood

- ▶ Training objective: Maximize marginal log-likelihood (MLL)

$$L(\theta) = \mathbf{E}_{x \sim \mathbf{pop}} [\log m_{\theta}(x)]$$

(Equivalent to cross entropy minimization, but convenient to frame as maximization for later)

- ▶ Requires the ability to calculate marginal probability of x !

$$m_{\theta}(x) = \mathbf{E}_{z \sim \pi_{\theta}} [\kappa_{\theta}(x|z)]$$

- ▶ Sometimes we can calculate it exactly (best scenario)
 - ▶ z is discrete and \mathcal{Z} is small: $m_{\theta}(x) = \sum_{z \in \mathcal{Z}} p_{\theta}(x, z)$ directly computable
 - ▶ p_{θ} makes Markov assumptions: $m_{\theta}(x)$ computable by dynamic programming (e.g., forward algorithm for HMMs)
- ▶ In general, we need to approximate by sampling